

JOHN DOE

📍 San Francisco, CA | ✉ john.doe@email.com | 🌐 rendercv.com | 🌐 rendercv | 🌐 rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and link.

Education

Princeton University

PhD — Computer Science

Princeton, NJ
Sept 2018 - May 2023

- Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment
- Advisor: Prof. Sanjeev Arora
- NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)

Boğaziçi University

BS — Computer Engineering

Istanbul, Türkiye
Sept 2014 - June 2018

- GPA: 3.97/4.00, Valedictorian
- Fulbright Scholarship recipient for Graduate Studies

Experience

Nexus AI

Co-Founder & CTO

June 2023 - present
San Francisco, CA

- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
- Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
- Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
- Developed proprietary inference optimization reducing latency by 73% compared to baseline

NVIDIA Research

Research Intern

May 2022 - Aug 2022
Santa Clara, CA

- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
- Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)

Google DeepMind

Research Intern

May 2021 - Aug 2021
London, UK

- Developed reinforcement learning algorithms for multi-agent coordination
- Published research at top-tier venues with significant academic impact - ICML 2022 main conference paper, cited 340+ times within two years - NeurIPS 2022 workshop paper on emergent communication protocols - Invited journal extension in JMLR (2023)

Apple ML Research

Research Intern

May 2020 - Aug 2020
Cupertino, CA

- Created on-device neural network compression pipeline deployed across 50M+ devices
- Filed 2 patents on efficient model quantization techniques for edge inference

- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

Projects

FlashInfer Open-source library for high-performance LLM inference kernels	Jan 2023 - present
• Achieved 2.8x speedup over baseline attention implementations on A100 GPUs • Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors	
NeuralPrune Automated neural network pruning toolkit with differentiable masks	Jan 2021
• Reduced model size by 90% with less than 1% accuracy degradation on ImageNet • Featured in PyTorch ecosystem tools, 4,200+ GitHub stars	

Publications

John Doe, Sarah Williams, David Park Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models — <i>NeurIPS 2023</i>	July 2023
James Liu, John Doe Neural Architecture Search via Differentiable Pruning — <i>NeurIPS 2022, Spotlight</i>	Dec 2022
Maria Garcia, John Doe, Tom Anderson Multi-Agent Reinforcement Learning with Emergent Communication — <i>ICML 2022</i>	July 2022
John Doe, Kevin Wu On-Device Model Compression via Learned Quantization — <i>ICLR 2021, Best Paper Award</i>	May 2021

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

- **Languages:** Python, C++, CUDA, Rust, Julia
- **ML Frameworks:** PyTorch, JAX, TensorFlow, Triton, ONNX
- **Infrastructure:** Kubernetes, Ray, distributed training, AWS, GCP
- **Research Areas:** Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)